

Analytical Version

Maxime Rischard

February 19, 2016

```
In [ ]: using Distributions
        using PDMats
        using Optim
        using StatsFuns: log2π
        import PyPlot; plt=PyPlot
        using DualNumbers
        using Gadfly
        using ToeplitzMatrices
        plt.svg(true)
        using Mamba
        using GraphViz
        using ProfileView
        using GaussianProcesses
        using ProfileView
        ;

In [ ]: # function sqexp{S<:Number,T<:Number}(σ_f::S, L::T)
        #     function k{T<:Number}(x::T, xprime::T)
        #         return σ_f^2 * exp( - (x-xprime)^2 / (2*L^2))
        #     end
        #     return k
        # end

        function kernel{T<:Number,Kern<:Kernel}(kern::Kern, x::AbstractVector{T})
            n = length(x)
            Xmat = reshape(x, (1,n))
            Σ = GaussianProcesses.crossKern(Xmat, kern)
            return Σ
        end
```

Sample 200 points uniformly on either side of the discontinuity at $x = 5$.

$$X_1 \sim \text{Uniform}(0, 5) \tag{1}$$

$$X_2 \sim \text{Uniform}(5, 10) \tag{2}$$

$$\tag{3}$$

```
In [ ]: _X1distr = Uniform(0,5)
        _X2distr = Uniform(5,10)
        _n_per_region = 200
        _n = _n_per_region*2
        _thresh = 5.0
        _X = sort(Float64[rand(_X1distr, _n_per_region); rand(_X2distr, _n_per_region)])
        ;
```

We want to simulate a GP; we'll simulate from the model

$$y = \tau \mathbf{I}\{x > 5\} + X\beta + \epsilon$$

X will be the locations, and y will be the response surface. The shift term $\tau \mathbf{I}\{x > 5\}$ will be zero to the left of the boundary, and a constant to the right of the boundary. ϵ is from a GP with a squared exponential covariance with additional iid normal noise with variance σ_y^2 .

$$K(x, x') = \text{cov}(y, y') = \sigma_y^2 \delta(x - x') + \sigma_{GP}^2 \exp\left(-\frac{(x - x')}{2L^2}\right)$$

Set the data-generating parameters to the following (arbitrary) values:

$$\sigma_f^2 = 1 \tag{4}$$

$$L^* = 1.05 \tag{5}$$

$$\tau^* = 0.75 \tag{6}$$

$$\beta^* = 1.0 \tag{7}$$

$$\sigma_y^{2*} = 0.01 \tag{8}$$

$$\tag{9}$$

```
In [ ]: # data-generating parameters
        _sf2_star = 1.0
        _Lstar   = 1.05
        _tstar   = 0.75
        _bstar   = 1.0
        _sy2_star = 0.01
        _region = _X.>_thresh
        ;
```

We are fitting a Bayesian model with normal priors on β and τ . Let's start by giving them very diffuse normal priors. The normality allows us to use all the tools of multivariate normal theory.

```
In [ ]: function simulate1(T<:Number,S<:Number,Ker<:Kernel)(beta::T, tau::T, kern::Ker, sy2::T, thresh::Real)
        region = X.>thresh
        n = length(X)
        K = kernel(kern, X) + 1e-5*eye(n)
        fXdistr = MvNormal(zeros(n), K)
        fX = rand(fXdistr)
        In=eye(Diagonal{Float64},n)
        Z_Ydistr = Normal(0, sqrt(sy2))
        Y = fX .+ beta*X + tau*region .+ rand(Z_Ydistr, n)
        return Y
    end
```

```
In [ ]: _sigma2 = 1000.0 # extremely diffuse normal prior on beta
        _sigma_tau2 = 1000.0 # extremely diffuse normal prior on tau
        _kern = SE(log(_Lstar), log(sqrt(_sf2_star)))
        _II = _region * _region'
        _Y = simulate1(_bstar, _tstar, _kern, _sy2_star, _thresh, _X)
        plt.plot(_X, _Y, "o")
        plt.title("Simulated Dataset")
        plt.xlabel("X")
        plt.ylabel("Y")
        ;
```

Conditionally on σ_y^2 , σ_f^2 and L , the posterior of τ (i.e. $\tau \mid Y, \sigma_y^2, \sigma_f^2, L$) can be obtain analytically. We have:

$$\tau \sim \mathcal{N}(0, \sigma_\tau^2) \quad (10)$$

$$\beta \sim \mathcal{N}(0, \sigma_\beta^2) \quad (11)$$

$$Y = \tau \mathbf{I}\{x > 5\} + X\beta + \underbrace{GP(0, K)}_f + \epsilon_{iid} \quad (12)$$

$$\text{cov}(Y, \tau) = \sigma_\tau \begin{pmatrix} 0 \\ \vdots \\ 0 \\ 1 \\ \vdots \\ 1 \end{pmatrix} \quad (13)$$

$$\text{cov}(Y, \beta) = \sigma_\beta X \quad (14)$$

$$\text{cov}(Y, \tau) \text{cov}(Y, \tau)^T = \sigma_\tau^2 \mathbf{I}\{x > 5\} \mathbf{I}\{x > 5\}^T = \sigma_\tau^2 \begin{bmatrix} 0 & 0 & \dots & 0 & 0 \\ 0 & 0 & \dots & 0 & 0 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \dots & 1 & 1 \\ 0 & 0 & \dots & 1 & 1 \end{bmatrix} \quad (15)$$

$$\text{cov}(Y, \beta) \text{cov}(Y, \beta)^T = \sigma_\beta^2 X X^T \quad (16)$$

$$\text{var}(Y) = \text{cov}(Y, \tau) \text{cov}(Y, \tau)^T + \text{cov}(Y, \beta) \text{cov}(Y, \beta)^T + K + \sigma_y^2 I \quad (17)$$

$$Y \sim \mathcal{N}(0, \text{var}(Y)) \quad (18)$$

$$(19)$$

$$\begin{pmatrix} Y \\ f \\ \tau \\ \beta \end{pmatrix} \sim \mathcal{N} \left(\begin{pmatrix} 0 \\ 0 \\ 0 \\ 0 \end{pmatrix}, \begin{bmatrix} \text{var}(Y) & K & \text{cov}(Y, \tau) & \text{cov}(Y, \beta) \\ K & K & 0 & 0 \\ \text{cov}(Y, \tau)^T & 0 & \sigma_\tau^2 & 0 \\ \text{cov}(Y, \beta)^T & 0 & 0 & \sigma_\beta^2 \end{bmatrix} \right)$$

From this we can obtain the posterior on τ :

$$\tau \mid Y \sim \mathcal{N} \left(\text{cov}(Y, \tau)^T \text{var}(Y)^{-1} Y, \sigma_\tau^2 - \text{cov}(Y, \tau)^T \text{var}(Y)^{-1} \text{cov}(Y, \tau) \right)$$

```
In [ ]: _XXt = _X*_X'
# marginal variance of Y
_VarY = kernel(_kern, _X) + _sigma_tau2*_II + _sigma_beta2*_XXt + _sigma_y2_star*eye(Diagonal{Float64}, _n)
_YDistr = MvNormal(Distributions.ZeroVector{Float64}, _n), PDMat(_VarY)
_VarYinv = inv(_VarY)
_Etau_Y = (_sigma_tau2 * _region' * _VarYinv * _Y)[1]
_Vtau_Y = (_sigma_tau2 - _sigma_tau2^2 * _region' * _VarYinv * _region)[1]
_tau_Y = Normal(_Etau_Y, sqrt(_Vtau_Y))

In [ ]: function mvnorm_logpdf{T<:Real, S<:Number}(\Sigma::Array{S,2}, X::Vector{T})
    \Sigmachol = cholfact(\Sigma)
    c0 = -0.5 * (length(X) * Float64(log2pi) + logdet(\Sigmachol))
    sqmahal = dot(X, inv(\Sigmachol) * X)
    return c0 - 0.5*sqmahal
end
```

We can also get the marginal likelihood of $\mathbb{P}(Y \mid \sigma_y^2, \sigma_f^2, L)$

```
In [ ]: mvnorm_logpdf(_VarY, _Y), logpdf(_YDistr, _Y)
```

```
In [ ]: @time mvnorm_logpdf(_VarY, _Y);
```

```
In [ ]: @time YDistr = MvNormal(Distributions.ZeroVector{Float64, _n}, PDMat(_VarY)); logpdf(_YDistr, _Y)
```

Maybe not such a great idea.

1 Partial Bayes

These computations are pretty quick, so we can optimize over $\sigma_y^2, \sigma_f^2, L$ (is this partial Bayes?).

```
In [ ]: """
Marginal variance of Y for fixed 'στ2', 'σβ2' and 'L'
"""

function g_VarY{T<:Number, S<:Number}(X::Vector{T}, XXt::Symmetric{T}, thresh::T, στ2::T, σβ2::T)
    n = length(X)
    kern = SE(log(L), log(√σf2))
    In=eye(Diagonal{Float64},n)
    region = X.>thresh
    II = region * region'
    K = kernel(kern, X)
    VarY = K + στ2*II + σβ2*XXt + σy2*In + 1e-5*eye(n)
    return VarY
end

function gen_optim_target{T<:Real}(X::Vector{T}, Y::Vector{T}, thresh::T, στ2::T, σβ2::T)
    XXt = Symmetric(X * X') # = BLAS.syrk('U', 'N', 1.0, X)
    function optim_target{T<:Number}(logx::Vector{T})
        σy2, σf2, L = exp(logx)
        VarY = g_VarY(X, XXt, thresh, στ2, σβ2, σy2, σf2, L)
        loglik = mvnorm_logpdf(VarY, Y)
        return -loglik
    end
end

optim_t = gen_optim_target(_X, _Y, 5.0, _στ2, _σβ2)
```

```
In [ ]: isapprox(g_VarY(_X, Symmetric(_XXt), 5.0, _στ2, _σβ2, _σy2_star, _σf2_star, _Lstar), _VarY)
```

```
In [ ]: @time _partial_opt=optimize(optim_t, [0.0, 0.0, 0.0], method=:l_bfgs; ftol=0.01, show_every=true)
```

It takes between 1 and 2 minutes with 200 points on each side of the discontinuity! And it shouldn't be any slower to use the same algorithm in higher dimensions.

```
In [ ]: function τ_posterior{T<:Real, S<:Number}(X::Vector{Float64}, Y::Vector{Float64}, στ2::T, σβ2::T)
    XXt = Symmetric(X * X')
    VarY = g_VarY(X, XXt, thresh, στ2, σβ2, σy2, σf2, L)
    VarYinv = inv(VarY)
    region = X.>thresh
    II = region * region' # we're doing this twice. expensive?
    Eτ_Y = (στ2 * region' * VarYinv * Y)[1]
    Vτ_Y = (στ2 - στ2^2 * region' * VarYinv * region)[1]
end
```

```

         $\tau_Y = \text{Normal}(E\tau_Y, \sqrt{V\tau_Y})$ 
    return  $\tau_Y$ 
end
;

```

Using these point estimates of the covariance parameters, and then computing the posterior on τ , we get a good estimate of τ (which hints that our estimator is consistent).

```

In [ ]: _ $\sigma y2$ _PB, _ $\sigma f2$ _PB, _L_PB = exp(_partial_opt.minimum)
         $\tau$ _posterior(_X, _Y, _ $\sigma \tau 2$ , _ $\sigma \beta 2$ , _ $\sigma y2$ _PB, _ $\sigma f2$ _PB, _L_PB, 5.0)

```

2 Slice Sampling

```

In [ ]: _ $\sigma f2$ _star, _ $\sigma f2$ _PB

```

```

In [ ]: _ $\sigma f2$ _prior = Truncated(Cauchy(0.0, 10.0), 0.0, Inf)
        _ $\sigma y2$ _prior = Truncated(Cauchy(0.0, 10.0), 0.0, Inf)
        _L_prior = Truncated(Cauchy(0.0, 10.0), 0.0, Inf)
        _xx = linspace(0.0,10.0,1000)
        plt.plot(_xx, pdf(_ $\sigma f2$ _prior, _xx))

```

```

In [ ]: _model = Model(
     $\sigma f2$  = Stochastic(
        () -> Truncated(Cauchy(0.0, 10.0), 0.0, Inf),
        true
    ),
     $\sigma y2$  = Stochastic(
        () -> Truncated(Cauchy(0.0, 10.0), 0.0, Inf),
        true
    ),
    L = Stochastic(
        () -> Truncated(Cauchy(0.0, 10.0), 0.0, Inf),
        true
    ),
    Y = Stochastic(1,
        (X, XXt, thresh,  $\sigma \tau 2$ ,  $\sigma \beta 2$ ,  $\sigma f2$ ,  $\sigma y2$ , L) -> MvNormal(
            Distributions.ZeroVector(Float64, n),
            PDMat(
                g_VarY(X, XXt, thresh,  $\sigma \tau 2$ ,  $\sigma \beta 2$ ,  $\sigma y2$ ,  $\sigma f2$ , L)
            )
        ),
        false
    ),
)

```

```

In [ ]: Graph(graph2dot(_model))

```

```

In [ ]: _nuts_scheme = [NUTS([: $\sigma f2$ , : $\sigma y2$ , :L])]
        _slice_scheme = [Slice([: $\sigma f2$ , : $\sigma y2$ , :L], 1.0)]
        _gibbsslice_scheme = [Slice([: $\sigma f2$ ], 1.0), Slice([: $\sigma y2$ ], 1.0), Slice([:L], 1.0)]
        _data = Dict{Symbol, Any}(
            :X => _X,
            :XXt => Symmetric(_X * _X'),
            :Y => _Y,
            :thresh => 5.0,

```

```

        : $\sigma\tau^2$  => _ $\sigma\tau^2$ ,
        : $\sigma\beta^2$  => _ $\sigma\beta^2$ ,
    )
    _pr = Truncated(Cauchy(0.0, 10.0), 0.0, Inf)
    ## Initial Values
    inits = [
        Dict{Symbol, Any}(
            :Y => _data[:Y],
            : $\sigma y^2$  => rand(_pr),
            : $\sigma f^2$  => rand(_pr),
            :L => rand(_pr),
        )
    ]
    for i in 1:3
    ]
;

In [ ]: setsamplers!(model, gibbsslice_scheme)
        sim1 = mcmc(model, data, inits, 1000, burnin=100, thin=2, chains=3)

In [ ]: plot(sim1)[1]

In [ ]: plot(sim1)[2]

In [ ]: plot(sim1)[3]

In [ ]: gelmandiag(sim1)

```

Not converging at all!

3 Toeplitz Structure

Instead of sampling X uniformly from 0 to 10, we will put X on a grid. That way, the Gaussian Process' covariance matrix is Toeplitz, and efficient algorithms exist to invert it.

```

In [ ]: function kernel{T<:Number,Kern<:GaussianProcesses.Stationary}(kern::Kern, x::LinSpace{T})
        vc = GaussianProcesses.crossKern(x', reshape([0.0,],(1,1)), kern)
        K_ST=SymmetricToeplitz(vec(vc))
        return K_ST
    end

In [ ]: _Xgrid = linspace(0,10,_n_per_region*2)
        _Ygrid = simulate1(_ $\beta$ star, _ $\tau$ star, _kern, _ $\sigma y^2$ _star, _thresh, _Xgrid)
        plt.plot(_Xgrid, _Ygrid, ".")
        plt.xlabel("X")
        plt.ylabel("Y")
        ;

In [ ]: _K_ST = kernel(_kern, _Xgrid)
        _K = kernel(_kern, collect(_Xgrid))
        ;

In [ ]: @assert isapprox(_K, _K_ST, atol=1e-4, norm=Base.LinAlg.vecnormInf)
        Base.LinAlg.vecnormInf(_K-_K_ST) # maximum distance

In [ ]: typeof(_K_ST)

```

```
In [ ]: inv(_K)
        @time inv(_K);
```

```
In [ ]: inv(_K_ST)
        @time inv(_K_ST);
```

Inversion using the Toeplitz structure is about ten times faster. The advantage goes up as the matrix gets larger.

```
In [ ]: function sherman_morrison{T<:Number}(Ainv::AbstractMatrix{T}, u::AbstractVector{T}, v::AbstractVector{T})
    vT = v'
    α::T = one(T) / (one(T)+(dot(v,Ainv*u)))
    return Ainv - α.*(Ainv*u)*(vT*Ainv)
end
```

```
In [ ]: # quick test
        @assert isapprox(inv(_K+eye(_n)+3.0*_X*_X'), sherman_morrison(inv(_K+eye(_n)), 3.0*_X,_X))
```

```
In [ ]: # function g_VarY{T<:Number, S<:Number}(X::Vector{T}, XXt::Symmetric{T}, thresh::T, στ2::T, σβ2::T)
#     n = length(X)
#     kern = SE(log(L), log(√σf2))
#     In=eye(Diagonal{Float64},n)
#     region = X.>thresh
#     II = region * region'
#     K = kernel(kern, X)
#     VarY = K + στ2*II + σβ2*XXt + σy2*In + 1e-5*eye(n)
#     return VarY
# end
```

```
function g_VarYInv{T<:Number}(ΣGP::SymmetricToeplitz{T}, X::AbstractVector{T}, thresh::T, στ2::T)
    vc = copy(ΣGP.vc)
    vc[1] += σy2+1e-5
    ΣGP_obs = SymmetricToeplitz(vc)
    ΣGP_inv = inv(ΣGP_obs)
#     ΣGP_inv = inv(full(ΣGP)+eye(n)*σy2)
    ΣGPβ_inv = sherman_morrison(ΣGP_inv, σβ2*X, X)
    region = 1.0.*(X.>thresh) # maybe converting a boolean to a float is stupid
    ΣGPβτ_inv = sherman_morrison(ΣGPβ_inv, στ2*region, region)
    ΣGPβτμ_inv = ΣGPβτ_inv
#     ΣGPβτμ_inv = sherman_morrison(ΣGPβ_inv, σμ2*ones(T,n), ones(T,n))
    return ΣGPβτμ_inv
end
```

```
In [ ]: g_VarYInv(_K_ST, collect(_Xgrid), _thresh, 2.0, 10.0, 1.5);
        @time g_VarYInv(_K_ST, collect(_Xgrid), _thresh, 2.0, 10.0, 1.5);
```

```
In [ ]: inv(g_VarY(collect(_Xgrid), Symmetric(_Xgrid*_Xgrid'), _thresh, _στ2, _σβ2, 1.5, _σf2_star, _Ls))
        @time inv(g_VarY(collect(_Xgrid), Symmetric(_Xgrid*_Xgrid'), _thresh, _στ2, _σβ2, 1.5, _σf2_star, _Ls))
```

```
In [ ]: # checking that the inverse is correct
        @assert isapprox(g_VarYInv(_K_ST, collect(_Xgrid), _thresh, 2.0, 10.0, 1.5)*g_VarY(collect(_Xgrid), Symmetric(_Xgrid*_Xgrid'), _thresh, _στ2, _σβ2, 1.5, _σf2_star, _Ls), ones(T,n))
```

That's not great. It seems to be doing the right thing, but with large numerical error.

3.1 Optimisation using Toeplitz

```

In [ ]: function mvnorm_logpdf{T<:Real, S<:Number}(Σ::AbstractMatrix{S}, Σinv::AbstractMatrix{S}, X::Vector{S})
    c0 = -0.5 * (length(X) * Float64(log2π) + logdet(Σ))
    sqmahal = dot(X, Σinv * X)
    return c0 - 0.5*sqmahal
end

function gen_optim_target{T<:Real}(X::LinSpace{T}, Y::Vector{T}, thresh::T, στ2::T, σβ2::T)
    XXt = Symmetric(X * X') # = BLAS.syrk('U', 'N', 1.0, X)
    function optim_target{T<:Number}(logx::Vector{T})
        σy2, σf2, L = exp(logx)
        kern = SE(log(L), log(σf2))
        K = kernel(kern, X)
        VarY = g_VarY(collect(X), XXt, thresh, στ2, σβ2, σy2, σf2, L)
        VarYInv = g_VarYInv(K, collect(X), thresh, στ2, σβ2, σy2);
        loglik = mvnorm_logpdf(VarY, VarYInv, Y)
        return -loglik
    end
end

In [ ]: _optim_ST = gen_optim_target(_Xgrid, _Ygrid, 5.0, 10.0, 10.0)

In [ ]: _partial_opt_TS=optimize(_optim_ST, [0.0, 0.0, 0.0], method=:l_bfgs; ftol=0.01, show_every=true)
Profile.clear()
Profile.init(n=10^7, delay=0.01)
@profile _partial_opt_TS=optimize(_optim_ST, [0.0, 0.0, 0.0], method=:l_bfgs; ftol=0.01, show_every=true)

In [ ]: Profile.print(combine=true, format=:flat)

In [ ]: _optim_t = gen_optim_target(collect(_Xgrid), _Ygrid, 5.0, 10.0, 10.0)
@time _partial_opt=optimize(_optim_t, [0.0, 0.0, 0.0], method=:l_bfgs; ftol=0.01, show_every=true)

In [ ]: _partial_opt.minimum

In [ ]: _partial_opt

In [ ]: _partial_opt_TS.minimum

In [ ]: _partial_opt_TS

In [ ]: _optim_ST(_partial_opt_TS.minimum)

In [ ]: _optim_ST(_partial_opt.minimum)

In [ ]: _optim_t(_partial_opt_TS.minimum)

In [ ]: _optim_t(_partial_opt.minimum)

```

4 Autodiff

Auto-differentiation using the DualNumbers package doesn't work well for multivariate analysis.

```

In [ ]: using ForwardDiff
    x = rand(3)
    h{T<:Number}(x::Vector{T}) = sum(sin, x) + prod(tan, x) * sum(sqrt, x);
    ∇h = ForwardDiff.gradient(h)
    ∇h(x)

```



```

In [ ]: x = [Dual(2.0,1.0), Dual(3.0,1.0), Dual(2.5, 1.0), Dual(1.3, 1.0)]
        XXt = x * x'
        cholfact(Symmetric(XXt))

In [ ]: realpart(XXt)

In [ ]: @time optim_t(x);
        ∇opt = ForwardDiff.gradient(optim_t)
        @time ∇opt(x)

In [ ]: using Base.LinAlg: chol, chol!
        # Linear algebra on dual matrices

        # Solve the matrix system
        #
        #   U'*M + M'*U = B
        #
        # for M, where B is symmetric and U is upper triangular. This looks a bit like
        # the *-Sylvester equation, but with a lot more structure. The solution
        # algorithm is basically a forward substitution method.
        function tri_ss_solve!(M, U, B)
            n = size(U,1)
            # Compute M row by row. The expression for the i'th row is broken into a
            # part involving matrix products of the previously computed rows of M with
            # parts of U, and a part involving the current row.
            for i = 1:n
                a = B[i,i:end]
                if i > 1
                    # Uses only parts of M computed in previous iterations.
                    a -= M[1:i-1,i]'*U[1:i-1,i:end] + U[1:i-1,i]'*M[1:i-1,i:end]
                end
                M[i,i] = a[1]/(2*U[i,i])
                println("Ui: ", U[i,i])
                println("Mi: ", M[i,i])
                println("Ui,*:", U[i,i+1:end])
                M[i,i+1:end] = (a[1,2:end] - M[i,i]*U[i,i+1:end]) / U[i,i]
            end
            return M
        end

        # Version of tri_ss_solve!() optimized for BLAS types
        # function tri_ss_solve!{T<:Base.LinAlg.BlasFloat}(M::AbstractMatrix{T},
        #                                           U::AbstractMatrix{T}, B::AbstractMatrix{T})
        #
        #     n = size(U,1)
        #     a = zeros(n)
        #     mi = zeros(n)
        #     ui = zeros(n)
        #     unit = one(T)
        #     for i = 1:n
        #         m = n-i+1
        #         # a[1:m] = B[i,i:n]
        #         for k=1:m
        #             a[k] = B[i,i+k-1]
        #         end
        #         if i > 1

```

```

#           # a -= M[1:i-1,i]'*U[1:i-1,i:end] + U[1:i-1,i]'*M[1:i-1,i:end]
#           for k=1:i-1
#               mi[k] = M[k,i]
#               ui[k] = U[k,i]
#           end
#           # Call BLAS directly to avoid temporary array copies.
#           BLAS.gemv!('T', -unit, sub(U,1:i-1,i:n), mi, unit, a)
#           BLAS.gemv!('T', -unit, sub(M,1:i-1,i:n), ui, unit, a)
#       end
#       M[i,i] = a[1]/(2*U[i,i])
#       # M[i,i+1:end] = (a[1,2:end] - M[i,i]*U[i,i+1:end]) / U[i,i]
#       for k=2:m
#           M[i,i+k-1] = (a[k] - M[i,i]*U[i,i+k-1]) / U[i,i]
#       end
#   end
#   return M
# end

```

```

# Cholesky factorization of arrays of dual numbers
#
# The returned matrix is upper triangular
function chol42{T}(A :: AbstractMatrix{Dual{T}})
    U = chol(real(A))
    B = epsilon(A)
    M = zeros(size(A))
    tri_ss_solve!(M, U, B)
    return dual(U,M)
end

```

```

In [ ]: Σhat = cov(rand(Normal(), (10000,5)))
        Σdual = Array{Dual{Float64}, size(Σhat)}
        for i in 1:5
            for j in 1:5
                Σdual[i,j] = Dual(Σhat[i,j], 1.0)
            end
        end
        chol42(Σdual)

```